

F5 AI Security Guardrails on Amazon EKS

This document captures the end-to-end deployment of F5 AI Security Guardrails on a single-node Amazon EKS cluster for a personal proof of concept.

The deployment follows the F5 AI Security Kubernetes Helm path and uses:

- Amazon EKS
- One GPU managed node group
- F5 AI Security Operator installed from Harbor
- Guardrails inference enabled
- Red Team disabled
- In-cluster PostgreSQL
- NGINX Ingress Controller
- Route 53 DNS
- Self-signed TLS certificate for lab access

This is a PoC build, not a production reference architecture. Secrets, license values, passwords, and registry credentials are intentionally excluded.

Final Environment

Item	Value
Public URL	https://guardrails.f5demo.io
AWS region	us-east-1
EKS cluster	f5-guardrails-poc
Kubernetes version	1.35
Managed node group	guardrails-gpu-ng
Instance type	g5.4xlarge
GPU	1 NVIDIA GPU
Node disk	150 GiB gp3
Database	In-cluster PostgreSQL
F5 operator namespace	f5-ai-sec
Moderator namespace	cai-moderator
Prefect namespace	prefect
Inference namespace	f5-ai-sec-inference
Ingress namespace	ingress-nginx

Architecture

```

Browser
|
| https://guardrails.f5demo.io
v
Route 53 CNAME
|
v
AWS Network Load Balancer
|
v
NGINX Ingress Controller
|
+-- /auth -> cai-moderator:8080 (Keycloak)
|
+-- /      -> cai-moderator:5500 (F5 AI Security UI/API)

```

Application namespaces:

```

f5-ai-sec
- F5 AI Security Operator

cai-moderator
- Moderator UI/API
- Embedded Keycloak auth
- PostgreSQL

prefect
- Prefect server
- Prefect worker
- Workflow registration job

f5-ai-sec-inference
- F5 inference service
- Guardrails scanner model

```

Important Sizing Note

The F5 runbook recommends more headroom for single-node deployments. For this PoC, `g5.4xlarge` was selected as the smallest practical AWS instance because it provides:

- 1 x NVIDIA A10G GPU with 24 GiB VRAM
- 16 vCPU
- 64 GiB memory

Observed Kubernetes allocatable resources were approximately:

```

nvidia.com/gpu: 1
cpu:            15890m
memory:         62117664Ki

```

If pods remain Pending, are OOMKilled, or model loading becomes unstable, resize to `g5.8xlarge`.

Prerequisites

Local tools:

```

aws --version
kubectl version --client=true
helm version --short
eksctl version

```

Required inputs:

- F5 AI Security license string
- Harbor registry username and password for `harbor.calypsoai.app`
- AWS CLI credentials with permission to create EKS, EC2, IAM, CloudFormation, ELB, and Route 53 resources

- Public hosted zone for the selected hostname
- LLM provider credentials for post-install application configuration

The PoC used:

```
Hostname: guardrails.f5demo.io
Database: in-cluster PostgreSQL
Product: Guardrails only
Red Team: disabled
TLS:      self-signed certificate
```

Preflight Checks

Verify AWS identity:

```
aws sts get-caller-identity
```

Confirm default region:

```
aws configure get region
```

List existing EKS clusters:

```
aws eks list-clusters --region us-east-1
```

Check GPU instance availability:

```
aws ec2 describe-instance-type-offerings \
  --region us-east-1 \
  --location-type availability-zone \
  --filters Name=instance-type,Values=g5.4xlarge
```

Check GPU quota:

```
aws service-quotas list-service-quotas \
  --service-code ec2 \
  --region us-east-1 \
  --query "Quotas[?contains(QuotaName, 'Running On-Demand G') || contains(QuotaName, 'Running On-Demand P')].[QuotaName,QuotaCode,Value]" \
  --output table
```

Verify Route 53 hosted zone:

```
aws route53 list-hosted-zones-by-name \
  --dns-name f5demo.io \
  --output json
```

Update eksctl

The installed `eksctl` version originally supported only Kubernetes versions up to `1.31`. AWS listed `1.35` as the default supported EKS version in `us-east-1`, so `eksctl` was upgraded before cluster creation.

```
brew upgrade eksctl
eksctl version
```

Create the EKS Cluster

Dry-run the cluster configuration first:

```
eksctl create cluster \
  --name f5-guardrails-poc \
  --region us-east-1 \
  --version 1.35 \
  --managed \
  --nodegroup-name guardrails-gpu-ng \
  --node-type g5.4xlarge \
  --nodes 1 \
  --nodes-min 1 \
  --nodes-max 1 \
  --node-volume-size 150 \
```

```
--node-ami-family AmazonLinux2023 \
--dry-run
```

Create the cluster:

```
eksctl create cluster \
  --name f5-guardrails-poc \
  --region us-east-1 \
  --version 1.35 \
  --managed \
  --nodegroup-name guardrails-gpu-ng \
  --node-type g5.4xlarge \
  --nodes 1 \
  --nodes-min 1 \
  --nodes-max 1 \
  --node-volume-size 150 \
  --node-ami-family AmazonLinux2023 \
  --install-nvidia-plugin \
  --vpc-nat-mode Disable
```

Notes:

- NAT gateway was disabled to reduce recurring PoC cost.
- The node uses public subnet egress.
- `eksctl` installed the NVIDIA device plugin automatically because the node group used an EKS optimized accelerated AMI.

Verify the node:

```
kubectl get nodes -o wide
kubectl get node -o jsonpath='{range .items[*]}{.metadata.name}{"\t"}{.status.allocatable.
nvidia\.com/gpu}{"\t"}{.status.allocatable.cpu}{"\t"}{.status.allocatable.memory}{"\n"}}{end}'
kubectl get pods -n kube-system -o wide
```

Expected result:

- One Ready node
- `nvidia-device-plugin-daemonset` Running
- `nvidia.com/gpu` allocatable value is 1

Prepare Harbor Pull Secret

Create the F5 namespace:

```
kubectl create namespace f5-ai-sec
```

Create the Harbor pull secret:

```
kubectl create secret docker-registry regcred \
  --docker-server='harbor.calypsoai.app' \
  --docker-username='<HARBOR_USERNAME>' \
  --docker-password='<HARBOR_PASSWORD>' \
  -n f5-ai-sec
```

Log in with Helm:

```
helm registry login harbor.calypsoai.app
```

Install the F5 AI Security Operator

```
helm upgrade --install f5-ai-security-operator \
  oci://harbor.calypsoai.app/calypsoai/f5-ai-security-operator-helm \
  --version 1.4.1 \
  -n f5-ai-sec
```

Verify:

```
helm list -n f5-ai-sec
kubectl get crd securityoperators.ai.security.f5.com
kubectl get pods -n f5-ai-sec -o wide
```

Expected:

```
controller-manager    1/1    Running
```

Apply Guardrails SecurityOperator

Create a Guardrails-only `SecurityOperator` custom resource.

The actual deployment used a generated PostgreSQL password and the F5 license from a local file. Do not commit the real license or generated passwords.

```
apiVersion: ai.security.f5.com/v1alpha1
kind: SecurityOperator
metadata:
  name: security-operator-demo
  namespace: f5-ai-sec
spec:
  registryAuth:
    existingSecret: "regcred"

  postgresql:
    enabled: true
    values:
      postgresql:
        auth:
          password: "<GENERATED_POSTGRES_PASSWORD>"

  jobManager:
    enabled: true

  moderator:
    enabled: true
    values:
      env:
        CAI_MODERATOR_BASE_URL: https://guardrails.f5demo.io
      secrets:
        CAI_MODERATOR_DB_ADMIN_PASSWORD: "<GENERATED_POSTGRES_PASSWORD>"
        CAI_MODERATOR_DEFAULT_LICENSE: "<F5_LICENSE_STRING>"

  inference:
    enabled: true
    values:
      inference:
        guardrails:
          enabled: true
        redteam:
          enabled: false
```

Apply it:

```
kubectl apply -f securityoperator.yaml
```

Verify namespace and pod creation:

```
kubectl get securityoperator -n f5-ai-sec
kubectl get ns
kubectl get pods -n f5-ai-sec
kubectl get pods -n cai-moderator
kubectl get pods -n prefect
kubectl get pods -n f5-ai-sec-inference
```

Storage Fix: Install EBS CSI

PostgreSQL initially stayed Pending because the PVC required `ebs.csi.aws.com`, but the EBS CSI driver was not installed.

The `gp2` StorageClass also was not marked as default.

Patch gp2 as default:

```
kubectl patch storageclass gp2 \
  -p '{"metadata":{"annotations":{"storageclass.kubernetes.io/is-default-class":"true"}}}'
```

Install the EKS Pod Identity Agent:

```
eksctl create addon \
  --cluster f5-guardrails-poc \
  --region us-east-1 \
  --name eks-pod-identity-agent \
  --force
```

Install the EBS CSI driver with the recommended pod identity association:

```
eksctl create addon \
  --cluster f5-guardrails-poc \
  --region us-east-1 \
  --name aws-ebs-csi-driver \
  --auto-apply-pod-identity-associations \
  --force
```

Verify:

```
aws eks describe-addon \
  --cluster-name f5-guardrails-poc \
  --addon-name aws-ebs-csi-driver \
  --region us-east-1 \
  --query 'addon.[status,health.issues]' \
  --output json

kubectl get pods -n kube-system \
  -l app.kubernetes.io/name=aws-ebs-csi-driver \
  -o wide
```

```
kubectl get pvc -n cai-moderator -o wide
```

Expected:

- EBS CSI add-on **ACTIVE**
- EBS CSI controller and node pods Running
- PostgreSQL PVC Bound

Prefect Workflow Registration Fix

The first workflow registration jobs failed because they started before the Prefect API was ready.

Verify Prefect API health from inside the cluster:

```
kubectl exec -n prefect deploy/prefect-server -- \
  /bin/sh -c 'python - <<PY
import urllib.request
print(urllib.request.urlopen("http://prefect-server.prefect.svc.cluster.local:4200/api/health",
timeout=5).read().decode()[:200])
PY'
```

Expected:

```
true
```

Delete failed early jobs:

```
kubectl delete job -n prefect \
  cai-workflows-1781048696 \
  cai-workflows-1781048699 \
  --ignore-not-found
```

Create a fresh workflow registration job:

```
kubectl create job -n prefect \
  cai-workflows-manual-20260610-0054 \
```

```
--from=cronjob/cai-workflows
```

Verify:

```
kubectll get pods -n prefect -o wide  
kubectll logs -n prefect job/cai-workflows-manual-20260610-0054 --tail=100
```

Expected:

- Workflow registration job Completed
- Prefect server Running
- Prefect worker Running

Install NGINX Ingress Controller

Add and update the Helm repository:

```
helm repo add ingress-nginx https://kubernetes.github.io/ingress-nginx  
helm repo update ingress-nginx
```

Install ingress-nginx:

```
helm upgrade --install ingress-nginx ingress-nginx/ingress-nginx \\  
  --namespace ingress-nginx \\  
  --create-namespace \\  
  --set controller.service.annotations.service\beta\kubernetes\io/aws-load-balancer-type=nlb \\  
  --set controller.service.annotations.service\beta\kubernetes\io/aws-load-balancer- \\  
  scheme=internet-facing \\  
  --set controller.service.externalTrafficPolicy=Local
```

Verify the AWS load balancer hostname:

```
kubectll get svc -n ingress-nginx ingress-nginx-controller -o wide
```

Create TLS Secret

For this PoC, a self-signed certificate was used.

```
tmpdir=$(mktemp -d /private/tmp/guardrails-tls.XXXXXX)
```

```
openssl req -x509 -nodes -days 365 -newkey rsa:2048 \\  
  -keyout "$tmpdir/tls.key" \\  
  -out "$tmpdir/tls.crt" \\  
  -subj "/CN=gardrails.f5demo.io" \\  
  -addext "subjectAltName=DNS:gardrails.f5demo.io"
```

```
kubectll create secret tls cai-moderator-tls \\  
  -n cai-moderator \\  
  --cert="$tmpdir/tls.crt" \\  
  --key="$tmpdir/tls.key" \\  
  --dry-run=client \\  
  -o yaml | kubectll apply -f -
```

```
rm -rf "$tmpdir"
```

Browser warning is expected because the certificate is self-signed.

Create Ingress

```
apiVersion: networking.k8s.io/v1  
kind: Ingress  
metadata:  
  name: cai-moderator-ingress  
  namespace: cai-moderator  
  annotations:  
    nginx.ingress.kubernetes.io/proxy-read-timeout: "3600"  
    nginx.ingress.kubernetes.io/proxy-send-timeout: "3600"  
    nginx.ingress.kubernetes.io/proxy-http-version: "1.1"
```

```

    nginx.ingress.kubernetes.io/ssl-redirect: "true"
spec:
  ingressClassName: nginx
  tls:
  - hosts:
    - guardrails.f5demo.io
      secretName: cai-moderator-tls
  rules:
  - host: guardrails.f5demo.io
    http:
      paths:
      - path: /auth
        pathType: Prefix
        backend:
          service:
            name: cai-moderator
            port:
              number: 8080
      - path: /
        pathType: Prefix
        backend:
          service:
            name: cai-moderator
            port:
              number: 5500

```

Apply:

```
kubectl apply -f cai-moderator-ingress.yaml
```

Verify:

```
kubectl get ingress -n cai-moderator -o wide
```

Create Route 53 Record

Create a CNAME from the public hostname to the ingress-nginx load balancer hostname.

```

aws route53 change-resource-record-sets \
  --hosted-zone-id <F5DEMO_IO_HOSTED_ZONE_ID> \
  --change-batch file://guardrails-cname.json

```

Example change batch:

```

{
  "Comment": "Point Guardrails PoC hostname to ingress-nginx load balancer",
  "Changes": [
    {
      "Action": "UPSERT",
      "ResourceRecordSet": {
        "Name": "guardrails.f5demo.io.",
        "Type": "CNAME",
        "TTL": 60,
        "ResourceRecords": [
          {
            "Value": "<INGRESS_NGINX_LOAD_BALANCER_HOSTNAME>"
          }
        ]
      }
    ]
  }
}

```

Verify the change:

```
aws route53 get-change --id <CHANGE_ID>
```

Expected:

```
INSYNC
```

Final Validation

Check the public UI route:

```
curl -k -I --max-time 20 https://guardrails.f5demo.io/
```

Expected:

HTTP/2 200

Check the auth route:

```
curl -k -I --max-time 20 https://guardrails.f5demo.io/auth/
```

Expected:

HTTP/2 302

location: https://guardrails.f5demo.io/auth/admin/

Final pod check:

```
kubectl get pods -n f5-ai-sec -o wide
kubectl get pods -n cai-moderator -o wide
kubectl get pods -n prefect -o wide
kubectl get pods -n f5-ai-sec-inference -o wide
```

Expected state:

f5-ai-sec:	
controller-manager	1/1 Running
cai-moderator:	
cai-moderator	1/1 Running
cai-moderator-postgres-cai-postgresql-0	1/1 Running
prefect:	
prefect-server	1/1 Running
prefect-worker	1/1 Running
cai-workflows-manual-*	Completed
f5-ai-sec-inference:	
f5-ai-sec-inference	1/1 Running
model-cai-phi-4-gptq-4bit	1/1 Running

First Login and Password Change

Initial credentials from the F5 runbook:

Username: admin
Password: pass

The main Guardrails UI account page does not expose password management.

Use the embedded Keycloak account console:

<https://guardrails.f5demo.io/auth/realms/calypsoai/account/#/security/signingin>

After login, change the password under the account security/sign-in area.

Useful Commands

Cluster:

```
kubectl get nodes -o wide
kubectl top nodes
kubectl get pods -A
```

F5 namespaces:

```
for ns in f5-ai-sec cai-moderator prefect f5-ai-sec-inference ingress-nginx; do
  echo "=== $ns ==="
  kubectl get pods -n "$ns"
done
```

Moderator logs:

```
kubectl logs -n cai-moderator deploy/cai-moderator --tail=100
```

Operator logs:

```
kubectl logs -n f5-ai-sec deploy/controller-manager --tail=100
```

Inference logs:

```
kubectl logs -n f5-ai-sec-inference deploy/f5-ai-sec-inference --tail=100
```

Prefect logs:

```
kubectl logs -n prefect deploy/prefect-server --tail=100
```

```
kubectl logs -n prefect deploy/prefect-worker --tail=100
```

Storage:

```
kubectl get pvc -n cai-moderator
```

```
kubectl describe pvc -n cai-moderator
```

Ingress:

```
kubectl get svc -n ingress-nginx ingress-nginx-controller -o wide
```

```
kubectl get ingress -n cai-moderator -o wide
```

Cleanup

To remove the PoC cluster and the AWS resources created by `eksctl`:

```
eksctl delete cluster \
  --name f5-guardrails-poc \
  --region us-east-1
```

Then remove the Route 53 CNAME if it is no longer needed.

Cost-Control Automation

The repository includes a reusable automation script for deleting the PoC when it is not needed and recreating it later with minimal interaction.

```
./scripts/guardrails-poc.sh up
./scripts/guardrails-poc.sh status
./scripts/guardrails-poc.sh down --yes
```

The script covers the full deployment path used in this PoC, including the issues encountered during the first manual deployment:

- EKS cluster creation with a single GPU node
- NVIDIA device plugin installation through `eksctl`
- EKS Pod Identity Agent installation
- AWS EBS CSI driver installation
- `gp2` StorageClass default annotation patch
- Harbor pull secret creation
- F5 AI Security Operator installation
- Guardrails-only `SecurityOperator` creation
- In-cluster PostgreSQL configuration
- Prefect workflow registration retry after Prefect API readiness
- NGINX Ingress Controller installation

- Self-signed TLS secret creation
- `/auth` and `/` ingress routing
- Route 53 CNAME upsert and delete
- Post-delete cleanup of available EBS volumes still tagged as owned by the PoC cluster
- Public URL validation

Configure Automation

Copy the example config:

```
cp config/guardrails-poc.env.example config/guardrails-poc.env
```

Edit `config/guardrails-poc.env` if you want to change cluster name, region, instance type, hostname, hosted zone, or secret file locations.

By default, the script references:

```
HARBOR_CREDENTIALS_FILE="/Users/y.elmashad/F5/Labs/ai/guardrail/license/harbor.txt"
F5_LICENSE_FILE="/Users/y.elmashad/F5/Labs/ai/guardrail/license/license.txt"
```

The Harbor credentials file format is:

```
<harbor username>
<harbor password>
```

The license file should contain the license string as its full content.

You can also override secrets at runtime without changing the config file:

```
export HARBOR_USERNAME='<new-harbor-username>'
export HARBOR_PASSWORD='<new-harbor-password>'
export F5_LICENSE_STRING='<new-license-string>'

./scripts/guardrails-poc.sh up
```

The script stores the generated PostgreSQL password in:

```
.secrets/postgres-password
```

That directory is ignored by Git. Remove the file if you want a fresh generated PostgreSQL password on the next deployment.

Test Without Touching AWS

Use dry-run mode:

```
./scripts/guardrails-poc.sh --dry-run up
./scripts/guardrails-poc.sh --dry-run status
./scripts/guardrails-poc.sh --dry-run down --yes
```

Start the PoC

```
./scripts/guardrails-poc.sh up
```

Expected runtime is roughly 25-45 minutes because the script creates a new EKS control plane, GPU node, add-ons, images, model pod, ingress load balancer, and DNS record.

Stop the PoC and Minimize Cost

```
./scripts/guardrails-poc.sh down --yes
```

This deletes:

- Route 53 `guardrails.f5demo.io` CNAME
- EKS cluster

- GPU node group
- AWS load balancer
- EBS volumes created through the cluster
- Any available leftover EBS volumes still tagged `kubernetes.io/cluster/f5-guardrails-poc=owned`
- CloudFormation stacks created by `eksctl`

This delete-and-recreate model is slower than scaling a node group to zero, but it removes the EKS control-plane and GPU-node cost instead of merely reducing compute cost.

After teardown, `status` should show a clean stopped state:

```
Guardrails PoC status
=====
Environment:      stopped
EKS cluster:      not found (f5-guardrails-poc in us-east-1)
DNS record:       deleted (guardrails.f5demo.io not present)
EBS volumes:      none found with cluster-owned tag
Kubernetes resources: skipped because the EKS cluster is deleted
Public URL:       offline (https://guardrails.f5demo.io)

Next actions:
  Start again:     ./scripts/guardrails-poc.sh up
  Re-check status: ./scripts/guardrails-poc.sh status
```

Production Considerations

For production, change the following:

- Use a trusted TLS certificate from ACM or cert-manager.
- Use external PostgreSQL such as Amazon RDS.
- Use private node networking and controlled egress.
- Enable backup and restore for PostgreSQL.
- Use larger or separate node groups for CPU and GPU workloads.
- Use `g5.xlarge` or larger for single-node stability.
- Avoid embedding secrets in shell history or static manifests.
- Store credentials in AWS Secrets Manager, External Secrets Operator, Sealed Secrets, or another approved secret manager.
- Configure monitoring, logging, and alerting.
- Validate license entitlements before enabling Red Team.

References

- F5 AI Security documentation: <https://docs.aisecurity.f5.com/>
- F5 AI Security Helm install guide: <https://docs.aisecurity.f5.com/get-started/get-started-helm.html>
- F5 AI Security API docs: <https://docs.aisecurity.f5.com/api-docs/>
- Amazon EKS documentation: <https://docs.aws.amazon.com/eks/>
- ingress-nginx Helm chart: <https://github.com/kubernetes/ingress-nginx>